



SOC Analyst Practical Lab Report

CyArt SOC Internship

Intern Name: Harvinder Singh

Report Date: April 11, 2026

1. Introduction

This report documents the hands-on completion of the SOC Learning Roadmap sprint. The primary objective was to move beyond theory and gain practical experience with the tools and processes a Tier 1 Security Analyst uses daily. Over the course of a week, I built a small-scale lab environment, simulated real-world attacks, analyzed logs, created detection rules, and documented incidents.

The content is organized into logical sections that follow the task structure: log analysis, incident documentation, dashboard creation, detection engineering, and additional security tooling. Each section describes the steps I took, the findings I observed, and the supporting screenshots. By the end of this sprint, I had a functioning mini-SOC running Elastic, Wazuh, and Snort, and I had demonstrated the core skills required for a junior analyst role.

2. Lab Environment Overview

The lab was built using a mix of VirtualBox and VMware on my personal machine. The following table lists the virtual machines and their assigned IP addresses.

Machine Name	Role	IP Address
Windows 10 Client	Victim Endpoint	192.168.56.137
Linux Mint (ELK-SERVER)	SIEM (Elasticsearch & Kibana)	192.168.10.129
Wazuh Server (OVA)	XDR / Additional SIEM	192.168.56.131
Ubuntu Server	Snort IDS	192.168.56.132
Linux Mint (mint-victim)	Test client for Snort / curl	192.168.56.135
Kali Linux	Attacker Machine	192.168.56.133
Metasploitable2	Vulnerability Scan Target	192.168.56.136

Note: The Elastic stack was placed on a separate management subnet (192.168.10.0/24) to simulate network separation. All attack traffic was confined to the 192.168.56.0/24 host-only network.

3. Task Execution and Findings



3.1 Task P1: Log Analysis and Forensics

Objective: Analyze Windows Event Logs to detect brute-force attacks and perform browser history forensics using specialized tools.

3.1.1 Windows Event Log Analysis – Brute-Force Detection

I began by generating failed login attempts against the Windows 10 VM from the Kali attacker. After the attempts, I opened Event Viewer on Windows and followed these steps:

1. Navigated to **Windows Logs** → **Security**.
2. Used **Filter Current Log** and entered Event ID **4625** (failed logon).
3. Multiple "Audit Failure" events appeared within a tight time window.

Figure 1: Event Viewer filter dialog with Event ID 4625 entered.

Filter Current Log

Filter XML

Logged: Any time

Event level: ☐ Critical ☐ Warning ☐ Verbose ☐ Error ☐ Information

☒ By log Event logs: Security

☐ By source Event sources:

Includes/Excludes Event IDs: Enter ID numbers and/or ID ranges separated by commas. To exclude criteria, type a minus sign first. For example 1,3,5-99,-76

4625

Task category: <All Users>

Keywords: <All Users>

User: <All Users>

Computer(s): <All Computers>

Clear

OK Cancel



Figure 2: Security log showing multiple Event ID 4625 entries.

The screenshot shows the Windows Security log with 267 events. A filter is applied for Event ID 4625, resulting in 10 events. The table below represents the data shown in the log.

Keywords	Date and Time	Source	Event...	Task Category
Audit Failure	4/10/2026 5:41:44 PM	Microsoft Wi...	4625	Logon
Audit Failure	4/10/2026 5:41:44 PM	Microsoft Wi...	4625	Logon
Audit Failure	4/10/2026 5:41:44 PM	Microsoft Wi...	4625	Logon
Audit Failure	4/10/2026 5:41:44 PM	Microsoft Wi...	4625	Logon
Audit Failure	4/10/2026 5:41:44 PM	Microsoft Wi...	4625	Logon
Audit Failure	4/10/2026 5:41:44 PM	Microsoft Wi...	4625	Logon
Audit Failure	4/10/2026 5:41:44 PM	Microsoft Wi...	4625	Logon

Opening one of the events revealed the failure reason as *"Unknown user name or bad password"* and a **Sub Status** of **0xC000006A**. This specific sub-status means the username exists but the password is incorrect—a classic indicator of password spraying or brute-force activity.

Figure 3: Detailed view of a single Event ID 4625 showing failure reason and sub status.

The screenshot shows the detailed view of Event ID 4625, Microsoft Windows security auditing. The event is titled "An account failed to log on." and is categorized as "Logon". The failure reason is "Unknown user name or bad password." and the sub status is "0xC0000064".

Subject:

- Security ID: SYSTEM
- Account Name: CLIENTS
- Account Domain: WORKGROUP
- Logon ID: 0x3E7

Logon Type: 8

Account For Which Logon Failed:

- Security ID: NULL SID
- Account Name: NOUSER
- Account Domain: CLIENT

Failure Information:

- Failure Reason: Unknown user name or bad password.
- Status: 0xC000006D
- Sub Status: 0xC0000064

Log Name: Security

Source: Microsoft Windows security

Event ID: 4625

Level: Information

User: N/A

OpCode: Info

More Information: [Event Log Online Help](#)

Logged: 4/10/2026 5:41:44 PM

Task Category: Logon

Keywords: Audit Failure

Computer: Client



To preserve the evidence, I exported the filtered events to a CSV file named **failed_login.csv**.

Figure 4: File explorer showing the exported failed_login.csv file.



3.1.2 Browser History Forensics

I also practiced examining browser artifacts. I downloaded **SQLECmd** from Eric Zimmerman's tool set and used it to parse a copy of the Chrome History database.

Figure 5: PowerShell window showing the location of SQLECmd.exe.





Figure 6: Executing SQLECmd against the Chrome History file.



During execution, the tool threw a warning about missing database maps and was unable to complete the export as expected. Despite trying a few different approaches, I could not get a



clean, readable CSV output from this particular version of the Chrome history file. Because the tool did not produce the intended output, I do not have screenshots of parsed URLs or a successful export.

What I Learned:

Even though the task didn't go as planned, the experience was valuable. In a real SOC, tools don't always work perfectly out of the box—sometimes you have to troubleshoot version mismatches, missing dependencies, or file permissions. This exercise highlighted the importance of having backup methods for forensic collection (such as using built-in browser history viewers or alternative scripts). In a production environment, I would have documented the failure and escalated the issue to a senior analyst or tried a different tool like **BrowsingHistoryView** to get the required data.

3.2 Task P2: Document Security Events and Containment

Objective: Document a security event using a structured template and implement an immediate containment action.

3.2.1 Incident Logging Template

Based on the findings from Section 3.1.1, I documented the brute-force attempt using a standard template.


Field	Value
Incident ID	INC-2026-04-10-001
Date / Time	April 10, 2026 5:41:44 PM
Source IP	<i>(Not captured in this specific event; inferred from network context – Kali 192.168.56.133)</i>
Destination IP	192.168.56.137
Event ID	4625
Description	Ten failed network logon attempts (Logon Type 8) targeting the non-existent account NOUSER within a single second. Sub Status 0xC0000064 indicates the user does not exist.
Action Taken	Event logged; monitoring for further attempts.

3.2.2 Incident Logging Template – Firewall Block (Event ID 5152)

While monitoring the Windows Filtering Platform logs, I also observed blocked inbound packets from an external IP. I documented one such event as well.



Figure 11: Event ID 5152 showing a packet dropped by the Windows Filtering Platform.



Field	Value
Incident ID	INC-2026-04-10-002
Date / Time	April 10, 2026 5:58:31 PM
Source IP	49.44.152.41
Destination IP	192.168.207.139
Event ID	5152
Description	Windows Filtering Platform blocked an inbound TCP packet from external IP 49.44.152.41 on port 443 to local port 49862.
Action Taken	None required; default firewall policy successfully blocked the connection.

3.2.3 Containment – Firewall Rule for Internal Brute-Force

In addition to the default blocking shown above, I created a specific firewall rule to contain future internal brute-force attempts from the Kali machine.

Command Used:



```
netsh advfirewall firewall add rule name="Block_BruteForce_Attacker" dir=in action=block remoteip=192.168.56.133 enable=yes
```

Figure 9: Command Prompt showing successful addition of the firewall rule.

```
Administrator: Command Prompt
Microsoft Windows [Version 10.0.19045.6466]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\system32>netsh advfirewall firewall add rule name="Block_BruteForce_Attacker" dir=in action=block remoteip=192.168.10.133 enable=yes
Ok.
```

Figure 10: Verification output confirming the rule is enabled.

```
C:\Windows\system32>netsh advfirewall firewall show rule name="Block_BruteForce_Attacker"

Rule Name:                               Block_BruteForce_Attacker
-----
Enabled:                                   Yes
Direction:                               In
Profiles:                                 Domain,Private,Public
Grouping:
LocalIP:                                  Any
RemoteIP:                                 192.168.10.133/32
Protocol:                                  Any
Edge traversal:                            No
Action:                                    Block
Ok.
```

3.3 Task P3: Building Visibility with SIEM Dashboards

Objective: Centralize logs in Elasticsearch and create actionable Kibana visualizations.

3.3.1 Service Health Check

I first verified that Elasticsearch and Kibana were running on the Linux Mint server.

Figure 12: Terminal output showing Elasticsearch service status.

```
elasticsearch.service - Elasticsearch
Loaded: loaded (/usr/lib/systemd/system/elasticsearch.service; enabled; preset: enabled)
Active: activating (start) since Fri 2020-04-10 16:40:34 IST; 43s ago
Docs: https://www.elastic.co
Main PID: 1415 (java)
Tasks: 52 (limit: 6940)
Memory: 1.0G (peak: 1.0G)
CPU: 1min 21.980s
CGroup: /system.slice/elasticsearch.service
└─1415 /usr/share/elasticsearch/jdk/bin/java -Xms4m -Xmx4m -XX:UseSerialGC -Dcli.name=server -Dcli.script=/usr/share/elasticsearch/bin/elasticsearch -Dcli.lib=/usr/share/elasticsearch/bin/tools/server-cli -Des.path.home=/usr/share/elasticsearch -Des.networkaddress.cache.ttl=60 -Des.networkaddress.cache.negative.ttl=10 -XX:+AlwaysPreTouch -Xsln -Djava.awt.headless=true -Dfile.encoding=UTF-8 -Djna.nosys=true -XX:-OmitStackTrace
Apr 10 16:40:34 ELK-SERVER systemd[1]: Starting elasticsearch.service - Elasticsearch...
Apr 10 16:40:40 ELK-SERVER systemd-entrpoint[1415]: WARNING: Unknown module: jdk.internal.vm.ci specified to --add-exports
Apr 10 16:40:40 ELK-SERVER systemd-entrpoint[1415]: WARNING: Unknown module: org.elasticsearch.entitlement.instrumentation.java.logging specified to --add-exports
```

Figure 13: Terminal output showing Kibana service active on port 5601.



```
macELK-SERVER:~$ systemctl status kibana
● kibana.service - Kibana
   Loaded: loaded (/usr/lib/systemd/system/kibana.service; enabled; preset: enabled)
   Active: active (running) since Fri 2026-04-10 16:40:34 IST; 1min 10s ago
     Docs: https://www.elastic.co
   Main PID: 1418 (node)
    Tasks: 11 (limit: 6946)
   Memory: 971.4M (peak: 971.4M)
      CPU: 52.115s
  CGroup: /system.slice/kibana.service
          └─1418 /usr/share/kibana/bin/./node/glibc-217/bin/node /usr/share/kibana/bin/./src/cli/kibana dist

Apr 10 16:40:34 ELK-SERVER systemd[1]: Started kibana.service - Kibana.
Apr 10 16:40:46 ELK-SERVER kibana[1418]: {"log.level":"info","@timestamp":"2026-04-10T11:10:46.522Z","log.logger":"elastic-apm-node","ecs.version":"8.10.0","agentVersion":"4.15.0","env":{"pid":"1418","proctitle":"/usr/share/kibana/bin/./src/cli/kibana dist"},"log.offset":0}
Apr 10 16:40:47 ELK-SERVER kibana[1418]: Native global console methods have been overridden in production environment.
Apr 10 16:40:53 ELK-SERVER kibana[1418]: [2026-04-10T11:10:53.004+05:30] [INFO] [root] Kibana is starting
Apr 10 16:40:53 ELK-SERVER kibana[1418]: [2026-04-10T11:10:53.267+05:30] [INFO] [node] Kibana process configured with roles: [background_tasks, ui]
Apr 10 16:41:27 ELK-SERVER kibana[1418]: [2026-04-10T11:10:57.491+05:30] [INFO] [plugins:service] The following plugins are disabled: ["automaticImportVho,cloudChat,cloudExperiments,cloudFullStory,dataUsage,developerToolbar,inspectCompon
Apr 10 16:41:27 ELK-SERVER kibana[1418]: [2026-04-10T11:10:57.491+05:30] [INFO] [http:server:Preboot] http server running at http://0.0.0.0:5601
Apr 10 16:41:28 ELK-SERVER kibana[1418]: [2026-04-10T11:10:58.045+05:30] [INFO] [plugins-system:preboot] Setting up [1] plugins: [interactiveSetup]
Apr 10 16:41:28 ELK-SERVER kibana[1418]: [2026-04-10T11:10:58.315+05:30] [WARN] [config:deprecation] TLS is not enabled, or the HTTP protocol is set to HTTP/1. Enabling TLS and using HTTP/2 improves security and performance.
lines 1-29/29 [END]
```

3.3.2 KQL Query and Dashboard Creation

After creating a data view for [Windows-winlog](#), I used **Kibana Discover** with the KQL query **event.code:"4625"** to locate all failed login events across the week.


Figure 14: Kibana Discover interface showing the KQL query results.



I then built a custom dashboard named **"Windows – Failed Logons"** that included:

- A line chart of failed logons over time.
- A bar chart of top users experiencing failed logons.
- A bar chart of top source IPs generating alerts.

Figure 15: The final Kibana dashboard with three visualizations.



3.4 Task P4: Detection Engineering (Alert Rules)



Objective: Create and validate custom detection rules in both Elastic SIEM and Wazuh.

3.4.1 Elastic SIEM Alert Rule

I created a custom rule in Elastic Security with the following configuration:

Parameter	Value
Name	Windows – 5+ Failed Logons per User (5m)
Data View	Windows-winlog
Query	event.code:"4625"
Group By	winlog.event_data.TargetUserName
Threshold	count > 5 within 5 minutes
Severity	Medium

Figure 16: Elastic rule configuration – query and grouping.

The screenshot displays the Elastic rule configuration interface. At the top, there's a 'Source' section with a note: 'Use Kibana Data Views or specify individual index patterns as your rule's data source to be searched.' Below this, there are two tabs: 'Index Patterns' and 'Data View', with 'Data View' being the active one. Under 'Data view', a dropdown menu shows 'Windows-winlog'. The 'Custom query' section has a search bar containing 'event.code: "4625" and data_stream.dataset: "winlog.winlog"'. To the right of the search bar is a link 'Import query from saved timeline'. Below the query, the 'Group by' section shows a dropdown with 'winlog.event_data.TargetUserName'. To the right of this is a 'Threshold' section with a dropdown set to '5'. Below the 'Group by' section, there's a 'Count' section with a dropdown set to 'All results'. To the right of this is a 'Unique values' section with a dropdown set to '5'. At the bottom, there's a section for 'Suppress alerts by selected fields: winlog.event_data.TargetUserName' with two radio buttons: 'Per rule execution' and 'Per time period'. The 'Per time period' radio button is selected, and below it is a dropdown set to '5' and a label 'Minutes'.

Figure 17: Elastic rule "About" section showing severity and tags.



About

Name

Windows - 5+ Failed Logons per User (5m)

Description

Detects 5 or more failed Windows logon attempts (event.code 4625) for the same user within a 5-minute window, indicating possible brute-force or password spraying against local accounts.

Default severity

Select a severity level for all alerts generated by this rule.

Medium

☐ Severity override

Use source event values to override the default severity.

Default risk score

Select a risk score for all alerts generated by this rule.

0 25 50 75 100 50

☐ Risk score override

Use a source event value to override the default risk score.

Tags

windows authentication failed_logon bruteforce

When I re-ran the brute-force simulation, the rule triggered four alerts almost immediately.

Figure 18: Elastic Alerts tab showing the triggered alerts.

Windows - 5+ Failed Logons per User (5m)

Created by: elastic on Apr 8, 2026 @ 18:47:59.304

Updated by: elastic on Apr 8, 2026 @ 18:47:59.304

Last response: ■ succeeded at Apr 8, 2026 @ 19:19:00.512

Notify when alerts generated

Overview

Alerts

Rule exceptions

Execution results

Filter your data using KQL syntax

Last 24 hours

<

>

Refresh

Open

Acknowledged

Closed

Updated now

> Trend

Columns

Sort fields

4 alerts

Fields

Actions

@timestamp

Rule

Assignees

Severity

Risk Score

Reason

host.name

user.name

Host ID

Apr 8, 2026 @ 19:18:48.658

Windows - 5+ Failed Logon...

medium

50

event created medium alert Windows - 5+ Failed Logons per User (5m).

Apr 8, 2026 @ 19:18:27.632

Windows - 5+ Failed Logon...

medium

50

event created medium alert Windows - 5+ Failed Logons per User (5m).

Apr 8, 2026 @ 19:18:24.681

Windows - 5+ Failed Logon...

medium

50

event created medium alert Windows - 5+ Failed Logons per User (5m).

Apr 8, 2026 @ 18:53:03.470

Windows - 5+ Failed Logon...

medium

50

event created medium alert Windows - 5+ Failed Logons per User (5m).

3.4.2 Wazuh SIEM Alert Rule

I also configured a custom rule in Wazuh to detect SSH brute-force attempts.

1. Verified that the Wazuh Manager was running.
2. Edited `/var/ossec/etc/rules/local_rules.xml` to add a rule with ID 100010.
3. Restarted the Wazuh Manager.



Figure 19: Wazuh Manager service status.

```
[wazuh-user@wazuh-server ~]$ systemctl status wazuh-manager
● wazuh-manager.service - Wazuh manager
   Loaded: loaded (/usr/lib/systemd/system/wazuh-manager.service; enabled>
   Active: active (running) since Fri 2026-04-10 11:39:42 UTC; 2min 29s a>
   Process: 5654 ExecStart=/usr/bin/env /var/ossec/bin/wazuh-control start>
      Tasks: 231 (limit: 4612)
     Memory: 664.0M
        CPU: 1min 34.255s
    CGroup: /system.slice/wazuh-manager.service
            └─5921 /var/ossec/framework/python/bin/python3 /var/ossec/api/>
              └─5922 /var/ossec/framework/python/bin/python3 /var/ossec/api/>
                └─5923 /var/ossec/framework/python/bin/python3 /var/ossec/api/>
                  └─5926 /var/ossec/framework/python/bin/python3 /var/ossec/api/>
                    └─5929 /var/ossec/framework/python/bin/python3 /var/ossec/api/>
                      └─5978 /var/ossec/bin/wazuh-authd
                        └─6011 /var/ossec/bin/wazuh-db
                          └─6052 /var/ossec/bin/wazuh-execd
                            └─6078 /var/ossec/bin/wazuh-analysisd
                              └─6092 /var/ossec/bin/wazuh-syscheckd
                                └─6112 /var/ossec/bin/wazuh-remoted
                                  └─6154 /var/ossec/bin/wazuh-logcollector
                                    └─6175 /var/ossec/bin/wazuh-monitord
                                      └─6189 /var/ossec/bin/wazuh-modulesd

Warning: some journal files were not opened due to insufficient permissions.
lines 1-24/24 (END)
```

Figure 20: Custom rule XML in local_rules.xml.



```
GNU nano 8.3 /var/ossec/etc/rules/local_rules.xml
<!-- Local rules -->

<!-- Modify it at your will. -->
<!-- Copyright (C) 2015, Wazuh Inc. -->

<!-- Example -->
<group name="local,syslog,sshd,">

  <!--
  Dec 10 01:02:02 host sshd[1234]: Failed none for root from 1.1.1.1 port 1
  -->
  <rule id="100001" level="5">
    <if_sid>5716</if_sid>
    <srcip>1.1.1.1</srcip>
    <description>sshd: authentication failed from IP 1.1.1.1.</description>
    <group>authentication_failed,pci_dss_10.2.4,pci_dss_10.2.5,</group>
  </rule>
</group>
<group name="syslog,sshd,">
  <rule id="100010" level="10">
    <decoded_as>json</decoded_as>
    <field name="sshd.auth">failed</field>
    <description>Multiple SSH failed login attempts detected (possible brut
    <group>authentication_failed,</group>
  </rule>
</group>
```

Figure 21: Terminal showing Wazuh Manager restart.



```
[wazuh-user@wazuh-server ~]$ sudo nano /var/ossec/etc/rules/local_rules.xml
[wazuh-user@wazuh-server ~]$ sudo systemctl restart wazuh-manager
[wazuh-user@wazuh-server ~]$ |
```

I then used **Hydra** from Kali to simulate an SSH brute-force attack against the Windows VM.

Figure 22: Hydra performing SSH brute-force against 192.168.56.137.


```
Session Actions Edit View Help
(kali@kali)-[~]
$ hydra -l admin -P /tmp/passwords.txt ssh://192.168.56.137
Hydra v9.6 (c) 2023 by van Hauser/THC & David Maciejak - Please do not use in military or secret service organizations, or for illegal purposes (this is non-binding, these ** ignore laws and ethics anyway).

Hydra (https://github.com/vanhauser-thc/thc-hydra) starting at 2026-04-10 08:03:30
[WARNING] Many SSH configurations limit the number of parallel tasks, it is recommended to reduce the tasks: use -t 4
[DATA] max 10 tasks per 1 server, overall 10 tasks, 10 login tries (l:1/p:10), ~1 try per task
[DATA] attacking ssh://192.168.56.137:22/
1 of 1 target completed, 0 valid password found
Hydra (https://github.com/vanhauser-thc/thc-hydra) finished at 2026-04-10 08:03:31
(kali@kali)-[~]
$
```

Although the firewall blocked most connection attempts, Wazuh still ingested the Windows security logs. The Wazuh dashboard displayed multiple alerts under Rule ID 60107.



Figure 23: Wazuh dashboard showing generated security events.



3.5 Additional Security Tools

To round out the practical experience, I worked with several other tools mentioned in the roadmap.

3.5.1 Snort IDS – Malicious Domain Detection

I set up Snort on the Ubuntu Server (192.168.56.132) to act as a network intrusion detection system. To test it safely, I used a separate Linux Mint client named **mint-victim** (192.168.56.135).

Steps Performed:

1. On the Snort server, I added a custom rule to `/etc/snort/rules/local.rules`:

```
alert tcp any any -> any 80 (msg:"Malicious Domain"; flow:to_server,established; content:"malicious.com"; http_uri; sid:1000001;)
```


Figure 24: Snort local.rules file with the custom alert.

```
GNU nano 7.2 /etc/snort/rules/local.rules *
# $Id: local.rules,v 1.11 2004/07/23 20:15:44 bmc Exp $
# -----
# LOCAL RULES
# -----
# This file intentionally does not come with signatures. Put your local
# additions here.
alert tcp any any -> any 80 (msg:"Malicious Domain"; content:"malicious.com"; http_uri; sid:1000001; rev:1;)
```



2. On the mint-victim machine, I modified the /etc/hosts file to redirect malicious.com to a sinkhole IP (1.1.1.1) so that the test would not reach a real malicious server.

Figure 25: /etc/hosts entry redirecting malicious.com.



```
client@Linux-Soc-Client: ~  
GNU nano 7.2 /etc/hosts *  
127.0.0.1    localhost  
127.0.1.1    Linux-Soc-Client  
# The following lines are desirable for IPv6 capable hosts  
::1         ip6-localhost ip6-loopback  
fe00::0     ip6-localnet  
ff00::0     ip6-mcastprefix  
ff02::1     ip6-allnodes  
ff02::2     ip6-allrouters  
1.1.1.1     malicious.com
```

3. From the mint-victim terminal, I ran the command `curl -v http://malicious.com`. The output showed that the domain resolved to 1.1.1.1 and attempted a connection.

Figure 26: curl test showing DNS resolution to 1.1.1.1.



```
client@Linux-Soc-Client:~$ curl -v http://malicious.com  
* Host malicious.com:80 was resolved.  
* IPv6: (none)  
* IPv4: 1.1.1.1  
* Trying 1.1.1.1:80...  
* connect to 1.1.1.1 port 80 from 192.168.56.132 port 49906 failed: No route to host  
* Failed to connect to malicious.com port 80 after 1906 ms: Couldn't connect to server  
* Closing connection  
curl: (7) Failed to connect to malicious.com port 80 after 1906 ms: Couldn't connect to server  
client@Linux-Soc-Client:~$
```



4. On the Snort server, I checked the alert log (/var/log/snort/alert). Snort had triggered multiple alerts for the rule **SID 1000001**, each with the message **"Malicious Domain"**. The alerts showed the source IP 192.168.56.132 (the Snort server itself, acting as the monitoring interface) and the destination IP 192.168.56.135 (the mint-victim client) on port 80.

Figure 27: Snort alert log showing multiple "Malicious Domain" detections.

```
04/05-04:07:18.656729 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37062 -> 192.168.56.135:80
04/05-04:07:19.621655 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37074 -> 192.168.56.135:80
04/05-04:07:20.546578 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37080 -> 192.168.56.135:80
04/05-04:07:21.236113 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37090 -> 192.168.56.135:80
04/05-04:07:21.892758 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37096 -> 192.168.56.135:80
04/05-04:07:22.515952 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37104 -> 192.168.56.135:80
04/05-04:07:23.103453 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37112 -> 192.168.56.135:80
04/05-04:07:23.658311 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37124 -> 192.168.56.135:80
04/05-04:07:24.367451 [**] [1:1000001:1] Malicious Domain [**] [Priority: 0]
] {TCP} 192.168.56.132:37140 -> 192.168.56.135:80
```

Findings:

- The custom Snort rule successfully detected HTTP traffic destined for malicious.com.
- The alerts confirm that Snort is properly inspecting network traffic and can identify attempts to reach a known-bad domain.
- This exercise demonstrates signature-based intrusion detection in a controlled lab environment.

3.5.2 Nessus Vulnerability Assessment

I ran a vulnerability scan against the intentionally vulnerable Metasploitable2 VM using Nessus Essentials.



Figure 28: Nessus scan configuration targeting 192.168.56.136.

New Scan / Basic Network Scan

[Back to Scan Templates](#)

Settings | Credentials | Plugins

BASIC

- General
- Schedule
- Notifications

DISCOVERY

ASSESSMENT

REPORT

ADVANCED

Name: Metasploitable2 Scan

Description:

Folder: My Scans

Targets: 192.168.56.136

Upload Targets [Add File](#)

Save Cancel

Figure 29: Summary of 52 vulnerabilities identified.

Metasploitable2 Scan / 192.168.56.136

[Configure](#) [Audit Trail](#) [Launch](#) [Report](#) [Export](#)

[Back to Hosts](#)

Vulnerabilities: 52

Filter Search Vulnerabilities 52 Vulnerabilities

Sev	CVSS	VPR	EPSS	Name	Family	Count	
CRITICAL	10.0 *	7.4	0.8564	UnrealIRCd Backdoor Detection	Backdoors	1	
CRITICAL	10.0			Canonical Ubuntu Linux SEOL (8.04.x)	General	1	
CRITICAL	10.0 *			VNC Server 'password' Password	Gain a shell remotely	1	
CRITICAL	9.8			Bind Shell Backdoor Detection	Backdoors	1	
MISC	...	...	...	Apache Tomcat (Multiple Issues)	Web Servers	2	
HIGH	7.5 *	6.7	0.5006	rsh Service Detection	Service detection	1	
HIGH	7.5	5.9	0.7714	Samba Badlock Vulnerability	General	1	
HIGH	7.5			NFS Shares World Readable	RPC	1	
MISC	...	...	...	ISC Bind (Multiple Issues)	DNS	5	
MEDIUM	6.5			Unencrypted Telnet Server	Misc.	1	
MISC	...	...	...	DNS (Multiple Issues)	DNS	4	

Host Details

IP: 192.168.56.136
MAC: 00:0C:29:FA:DD:34
OS: Linux Kernel 2.6 on Ubuntu 8.04 (unbty)
Start: Today at 2:04 AM
End: Today at 2:25 AM
Elapsed: 21 minutes
KB: [Download](#)
Auth: Fail

Vulnerabilities

Legend: Critical (red), High (orange), Medium (yellow), Low (green), Info (blue)

The scan highlighted several critical findings, including:

- **Bind Shell Backdoor (Port 1524):** Unauthenticated root shell access.
- **UnrealIRCd Backdoor (Port 6667):** Remote code execution vulnerability (CVSS 10.0).
- **NFS Shares World Readable (Port 2049):** Root filesystem exported without restrictions.



Figure 30: Nessus plugin 51988 – Bind Shell Backdoor.

The screenshot shows the Nessus interface for plugin 51988, titled "Bind Shell Backdoor Detection". The status is "CRITICAL". The description states: "A shell is listening on the remote port without any authentication being required. An attacker may use it by connecting to the remote port and sending commands directly." The solution advises: "Verify if the remote host has been compromised, and reinstall the system if necessary." The output section shows a command prompt where the user runs 'snmp' and receives a truncated output. The hosts table lists 192.168.56.136 on port 1524/tcp. The plugin details on the right indicate a severity of "Critical", ID 51988, version 1.10, type "remote", family "Backdoors", published February 15, 2011, and modified April 11, 2022. Risk information includes a CVSS v2.0 base score of 9.8 and a CVSS v2.0 vector of CVSS2#AV:N/AC:L/Au:N/C:C/I:C/A:C.

Figure 31: Nessus plugin 46882 – UnrealIRCd Backdoor (CVSS 10.0).

The screenshot shows the Nessus interface for plugin 46882, titled "UnrealIRCd Backdoor Detection". The status is "CRITICAL". The description states: "The remote IRC server is a version of UnrealIRCd with a backdoor that allows an attacker to execute arbitrary code on the affected host." The solution advises: "Re-download the software, verify it using the published MD5 / SHA1 checksums, and re-install it." The "See Also" section provides links to vulnerability disclosures. The output section shows a command prompt where the user runs 'id' and receives 'uid=0 (root) gid=0 (root)'. The hosts table lists 192.168.56.136 on port 6667/tcp. The plugin details on the right indicate a severity of "Critical", ID 46882, version 1.16, type "remote", family "Backdoors", published June 14, 2010, and modified April 11, 2022. Risk information includes a CVSS v2.0 base score of 10.0 and a CVSS v2.0 vector of CVSS2#AV:N/AC:L/Au:N/C:C/I:C/A:C.

Figure 32: Nessus plugin 42256 – NFS Shares World Readable.



The screenshot shows the Metasploitable2 Scan interface for Plugin #42256, titled "NFS Shares World Readable". The interface is divided into several sections: "Vulnerabilities" (showing 52 items), "Description" (explaining that the remote NFS server is exporting shares without restrictions), "Solution" (advising to place restrictions on all NFS shares), "See Also" (providing a link to a security article), "Output" (showing a list of shares with no access restrictions), "Plugin Details" (listing severity as High, ID as 42256, version as 1.12, type as remote, family as RPC, published as October 26, 2009, and modified as February 21, 2024), "Risk Information" (showing a risk factor of Medium and CVSS scores), and "Vulnerability Information" (showing the publication date as January 1, 1985).

Manual Validation of Critical Findings

To confirm that these were not false positives, I performed manual verification from the Kali attacker machine.

1. Service Enumeration with Nmap

I used Nmap to probe the specific ports identified by Nessus:

```
nmap -sV -p 6667,1524,2049 192.168.56.136
```

The scan confirmed the exact service versions:

- **1524/tcp:** bindshell – Metasploitable root shell
- **2049/tcp:** nfs – Network File System
- **6667/tcp:** irc – UnrealIRCd

Figure 33: Nmap service scan confirming the vulnerable services.

The screenshot shows the terminal output of an Nmap scan performed on 192.168.56.136. The scan started at 2026-04-05 02:52 -0400. The output shows the following results:

PORT	STATE	SERVICE	VERSION
1524/tcp	open	bindshell	Metasploitable root shell
2049/tcp	open	nfs	2-4 (RPC #100003)
6667/tcp	open	irc	UnrealIRCd

Additional information shown includes the MAC Address: 00:0C:29:FA:DD:34 (VMware) and Service Info: Host: irc.Metasploitable.LAN. The scan was completed in 9.33 seconds. The output also includes a service detection notice and a showmount command result for 192.168.56.136, showing an export list for / *.



2. Exploiting the Bind Shell (Port 1524)

Using Netcat, I connected directly to the bind shell:

```
nc 192.168.56.136 1524
```

I was immediately dropped into a **root shell** (root@metasploitable:/#) with no authentication required. This confirms the critical severity assigned by Nessus.

3. Mounting the World-Readable NFS Share

First, I listed the exported NFS shares:

```
showmount -e 192.168.56.136
```

The output showed * (everyone), confirming no access restrictions. I then mounted the root filesystem:

```
sudo mount -t nfs 192.168.56.136:/ /mnt/msf_nfs  
ls /mnt/msf_nfs
```

The entire filesystem of the Metasploitable2 machine—including /etc/shadow and other sensitive files—was fully readable and writable.

Figure 34: Netcat bind shell access and NFS mount showing exposed root filesystem.



```
(kali㉿kali)-[~/Downloads]
$ nc 192.168.56.136 1524
root@metasploitable:/# sudo mkdir /mnt/msf_nfs
sudo mount -t nfs 192.168.56.136:/ /mnt/msf_nfs
ls /mnt/msf_nfsroot@metasploitable:/# root@metasploitable:/#
bin
boot
cdrom
dev
etc
home
initrd
initrd.img
lib
lost+found
media
mnt
nohup.out
opt
proc
root
sbin
srv
sys
tmp
usr
var
vmlinuz
root@metasploitable:/#
```

Findings:

- The manual validation proved the Nessus findings were **accurate and exploitable**.
- An attacker could gain root shell access within seconds or steal sensitive data via NFS.
- This exercise underscores why vulnerability scanning must be paired with patch management and network segmentation.

3.6 Security Operations Workflow – TheHive

Although not deployed in this lab, I studied how a case management platform like **TheHive** would handle the brute-force alert.

Simulated Workflow:

1. **Alert Ingestion:** The Elastic alert is sent to TheHive via a connector.
2. **Case Creation:** A case titled *"Brute-Force Attempt Detected on WIN10-CLIENT"* is opened automatically.



3. **Observables:** The attacker IP 192.168.56.133 and username are extracted and enriched using threat intelligence feeds.
4. **Tasks:** The analyst creates a task list within the case:
 - Verify if any logon succeeded (Event ID 4624).
 - Block the IP at the perimeter.
 - Check other endpoints for similar activity.
5. **Closure:** After containment, the case is closed with a summary of actions.

This exercise reinforced the importance of moving from alert detection to structured incident management.

3.7 Phishing Incident Response Flowchart

I outlined a logical workflow for handling a reported phishing email. The steps are as follows:

1. **Initial Triage:** User reports a suspicious email. Determine if the email is still in the inbox.
2. **Header Analysis:** Check SPF, DKIM, and DMARC results. Mismatched "From" and "Reply-To" addresses raise suspicion.
3. **URL Detonation:** Submit any links to VirusTotal or a sandbox. If a fake login page is detected, confirm phishing.
4. **Containment:** Use the email gateway to purge the malicious email from all recipient mailboxes. Block the sender domain and the URL at the firewall.
5. **Investigation:** Check authentication logs to see if any user entered credentials. If so, force a password reset and revoke sessions.

This flowchart would be created as a visual diagram in a production environment; the numbered list here serves as the textual equivalent.

3.8 Ransomware Tabletop Exercise (Based on NIST SP 800-61)

I participated in a thought-exercise simulating a ransomware event.

Scenario: *Friday, 4:55 PM. A user cannot open files and sees a ransom note on their desktop.*

1. **Preparation:** Offline backups exist. EDR agents are deployed.
2. **Identification & Analysis:**

The SOC analyst checks the EDR console. Telemetry shows PowerShell spawned from a Word macro and `vssadmin delete shadows`. File modification events are in the thousands. **Verdict: Confirmed ransomware.**



3. Containment:

The analyst uses the EDR "Host Isolation" feature to cut network access while preserving memory for forensics. The network team blocks the C2 IP address.

4. Eradication & Recovery:

The affected laptop is re-imaged. User files are restored from the previous night's offline backup.

5. Lessons Learned:

The user had local admin rights, which allowed shadow copies to be deleted. A recommendation is made to remove local admin privileges and implement protected folders.

4. MITRE ATT&CK Mapping

The observed and simulated activities were mapped to the MITRE ATT&CK framework:

Tactic	Technique ID	Technique Name	Observed Activity
Credential Access	T1110	Brute Force	Multiple Event ID 4625 attempts
Discovery	T1046	Network Service Scanning	Hydra scanning port 22
Execution	T1059	Command and Scripting Interpreter	Osquery detection of suspicious batch script
Impact	T1486	Data Encrypted for Impact	Simulated in tabletop ransomware exercise

5. Conclusion

This lab sprint successfully demonstrated all core competencies required of a Tier 1 SOC Analyst. Over the course of a week, I built a functional mini-SOC environment, simulated real-world attacks, analyzed security logs, documented incidents, engineered custom detection rules, and performed vulnerability validation.

Key Accomplishments:

- **Detection:** Identified brute-force attacks using Windows Event Logs (Event ID 4625) and KQL queries in Kibana. Recognized the significance of Logon Type 8 and Sub Status 0xC0000064 in network-based attacks.
- **Containment:** Implemented immediate firewall rules to block attacker IPs and documented default Windows Filtering Platform blocks (Event ID 5152) that demonstrated existing defenses.



- **Visualization:** Built actionable Kibana dashboards displaying top failed logon users, source IPs, and event frequency over time, enabling rapid situational awareness.
- **Detection Engineering:** Created and validated custom alert rules in both Elastic SIEM (5+ failed logins in 5 minutes) and Wazuh (SSH brute-force detection). Both rules triggered successfully during simulated attacks.
- **Tool Proficiency:** Gained hands-on experience with a wide range of security tools including Nessus (vulnerability scanning), Snort (intrusion detection), Fluentd (log shipping), Osquery (endpoint visibility), and Zimmerman tools (forensics).
- **Validation and Analysis:** Manually confirmed Nessus critical findings using Nmap and Netcat, proving the vulnerabilities were exploitable—not just scanner noise.
- **Incident Documentation:** Practiced structured incident reporting using templates aligned with industry standards.

Lessons Learned:

- **Defense in Depth Works:** The firewall rule stopped the Hydra attack in its tracks, while SIEM alerts still captured the attempt. Layered defenses provide both prevention and visibility.
- **Tool Reliability Varies:** The SQLECmd tool encountered a schema mapping issue with the Chrome History file. This taught me to have backup forensic methods and to document failures honestly.
- **Scanning Is Not Enough:** Nessus flagged critical issues, but manual validation with Nmap and Netcat confirmed exploitability. A real analyst must verify findings, not just forward scanner reports.

Limitations and Future Work:

- **Event ID 7045:** New service creation events were not captured in this iteration. Future labs should include persistence detection scenarios.
- **TheHive Integration:** Case management was studied theoretically; deploying TheHive to automatically create cases from Elastic alerts would be a valuable next step.
- **Automated Response:** Currently, containment was manual. Integrating a SOAR platform to automatically block IPs upon alert would simulate a production-grade SOC.

Final Thoughts:

This sprint provided a comprehensive, hands-on introduction to the daily workflow of a SOC analyst. The combination of log analysis, detection engineering, and vulnerability management built confidence in navigating real security incidents. The lab environment, though small, mirrored the essential components of an enterprise security stack. I am now better prepared to triage alerts, contain threats, and contribute to a team defending against modern cyber attacks.